
Deep Learning using TensorFlow

-By Prakash V

Intro

Introduction to Tensorflow and brief overview on architecture, trends and performance.

Neural Network using Numpy.

Neural Network using Tensorflow.

Deeper into Computer Vision

Fractional Max Pooling on MNIST Dataset.

Fractional Max Pooling on CIFAR-100 Dataset.

Trends - Deep Learning Libraries

Deep learning libraries: growth over past three months

new contributors from 2016-10-09 to 2017-02-10

#1:	192	tensorflow/tensorflow
#2:	89	dmlc/mxnet
#3:	78	fchollet/keras
#4:	42	baidu/paddle
#5:	29	Microsoft/CNTK
#6:	23	pfnet/chainer
#7:	21	Theano/Theano
#8:	20	deeplearning4j/deeplearning4j
#9:	20	tflearn/tflearn
#10:	19	BVLC/caffe
#11:	9	torch/torch7
#12:	3	NVIDIA/DIGITS

new forks from 2016-10-09 to 2017-02-10

#1:	6525	tensorflow/tensorflow
#2:	1822	BVLC/caffe
#3:	1316	fchollet/keras
#4:	999	dmlc/mxnet
#5:	909	deeplearning4j/deeplearning4j
#6:	887	Microsoft/CNTK
#7:	324	tflearn/tflearn
#8:	321	baidu/paddle
#9:	287	Theano/Theano
#10:	257	torch/torch7
#11:	175	NVIDIA/DIGITS
#12:	142	pfnet/chainer

new issues from 2016-10-09 to 2017-02-10

#1:	1563	tensorflow/tensorflow
#2:	979	fchollet/keras
#3:	871	dmlc/mxnet
#4:	646	baidu/paddle
#5:	486	Microsoft/CNTK
#6:	361	deeplearning4j/deeplearning4j
#7:	318	BVLC/caffe
#8:	217	NVIDIA/DIGITS
#9:	214	Theano/Theano
#10:	167	tflearn/tflearn
#11:	150	pfnet/chainer
#12:	90	torch/torch7

aggregate metrics growth from 2016-10-09 to 2017-02-10

#1:	54.01	tensorflow/tensorflow
#2:	18.71	fchollet/keras
#3:	16.38	dmlc/mxnet
#4:	12.86	BVLC/caffe
#5:	10.17	Microsoft/CNTK
#6:	9.32	baidu/paddle
#7:	8.75	deeplearning4j/deeplearning4j
#8:	4.21	Theano/Theano
#9:	3.89	tflearn/tflearn
#10:	3.14	NVIDIA/DIGITS
#11:	2.90	pfnet/chainer
#12:	2.46	torch/torch7

Neural networks

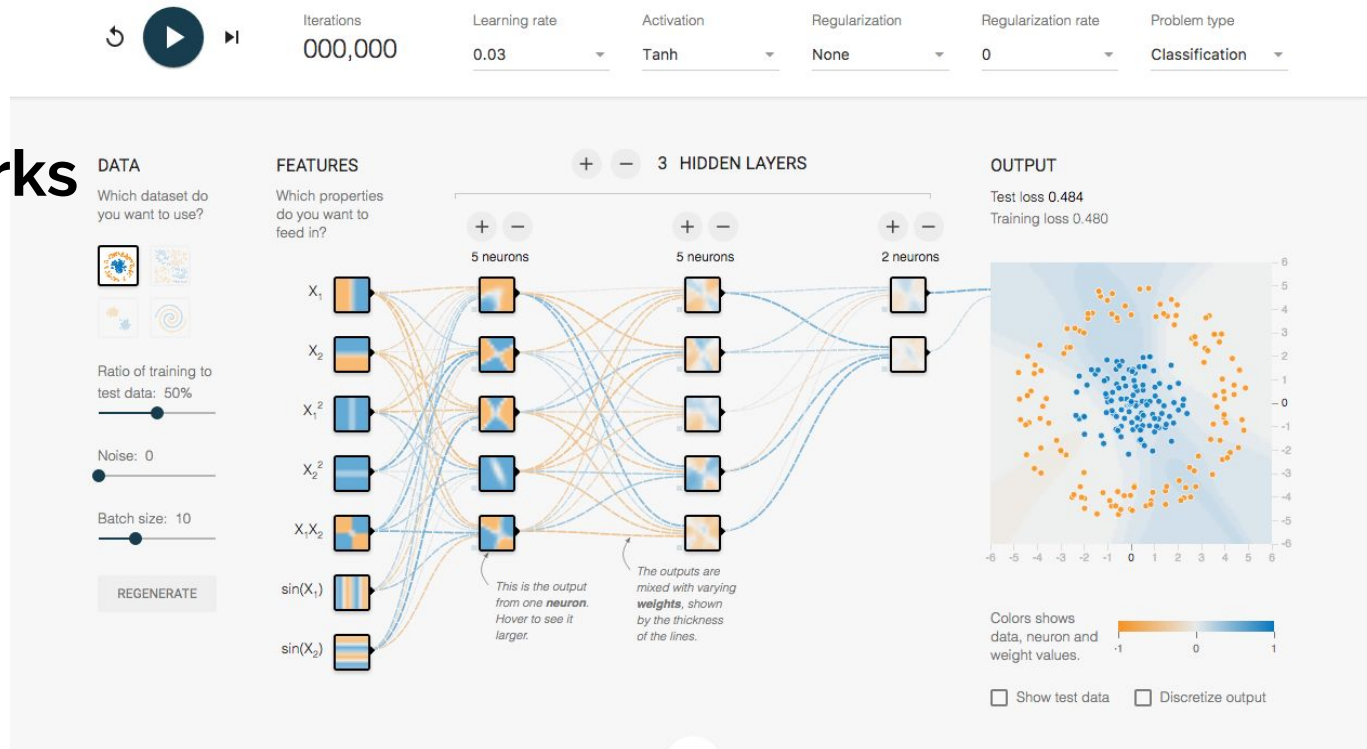
Input Data

Number of Layers

Activation Functions

Loss Function

Gradient optimizer



MNIST DATASET

Train Images - 60000

Test Images - 10000

Image size - 28 * 28 pixels

Class - 0, 1, 2, 3, 4, 5, 6, 7, 8, 9

Neural Network using Numpy

- A 3 layer Neural network.
- Initialize the weights randomly.
- Fix a learning rate.
- Use sigmoid as the activation function.

```
import numpy as np
from scipy import special

class neuralNetwork:
    #initialise the neural network
    def __init__(self, inputnodes, hiddennodes, outputnodes, learningrate):
        self.inodes = inputnodes
        self.hnodes = hiddennodes
        self.onodes = outputnodes

        self.wih = np.random.normal(0.0, pow(self.hnodes, -0.5), (self.hnodes, self.inodes))
        self.who = np.random.normal(0.0, pow(self.onodes, -0.5), (self.onodes, self.hnodes))

        self.lr = learningrate
        self.activation_function = lambda x: special.expit(x)

    pass
```

Neural Network using Numpy-2

- Give Inputs and Targets
- For each call to the train
 - A forward pass happens
 - Errors are calculated
 - Mean squared error is used for calculating the error.
 - Errors are back propagated.

```
def train(self, input_list, targets_list):
    inputs = np.array(input_list, ndmin=2).T
    targets = np.array(targets_list, ndmin=2).T

    #calculate signals into hidden layer
    hidden_inputs = np.dot(self.wih, inputs)
    #calculate the signals emerging from hidden layer
    hidden_outputs = self.activation_function(hidden_inputs)

    #calculate signals into final output layer
    final_inputs = np.dot(self.who, hidden_outputs)
    #calculate the signal emerging from final output layer
    final_outputs = self.activation_function(final_inputs)

    #output error
    output_errors = targets - final_outputs
    hidden_errors = np.dot(self.who.T, output_errors)

    #update the errors
    self.who += self.lr*np.dot((output_errors * final_outputs * (1.0 - final_outputs)), np.transpose(hidden_outputs))
    self.wih += self.lr * np.dot((hidden_errors* hidden_outputs*(1.0-hidden_outputs)), np.transpose(inputs))
```

Neural Network using Numpy-3

- How to test the result ?
- We write a query function to do

that.

```
#query the neural network
def query(self, input_list):
    inputs = np.array(input_list, ndmin=2).T
    hidden_inputs = np.dot(self.wih, inputs)
    hidden_outputs = self.activation_function(hidden_inputs)

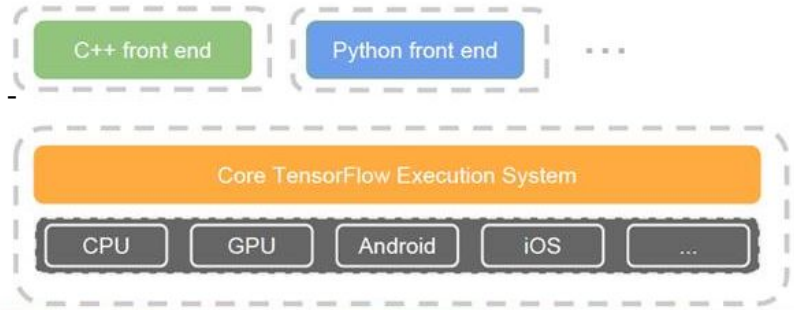
    #calculate signals into final output layer
    final_inputs = np.dot(self.who, hidden_outputs)
    final_outputs = self.activation_function(final_inputs)
    return final_outputs
```

Feed Forward Neural Networks using Numpy



About TensorFlow

- Open source software library for numerical computation using data flow graphs
- Nodes - operations , Edges - multi-dimensional Arrays (Tensors)
- Deploy computations to one or more CPUs and GPUs in Desktop, server, mobile device with single API
- Connect research and production - TensorFlow Serving.
- Auto Differentiation - Tensorflow handles computing the derivatives for you.
- Available in your favourite languages- Python , C++, Java, Go. Interfaces- Lua, JavaScript and R



Advantage of using Tensorflow

Less Overhead

Queues for input data

Multi-GPU, CPU
support

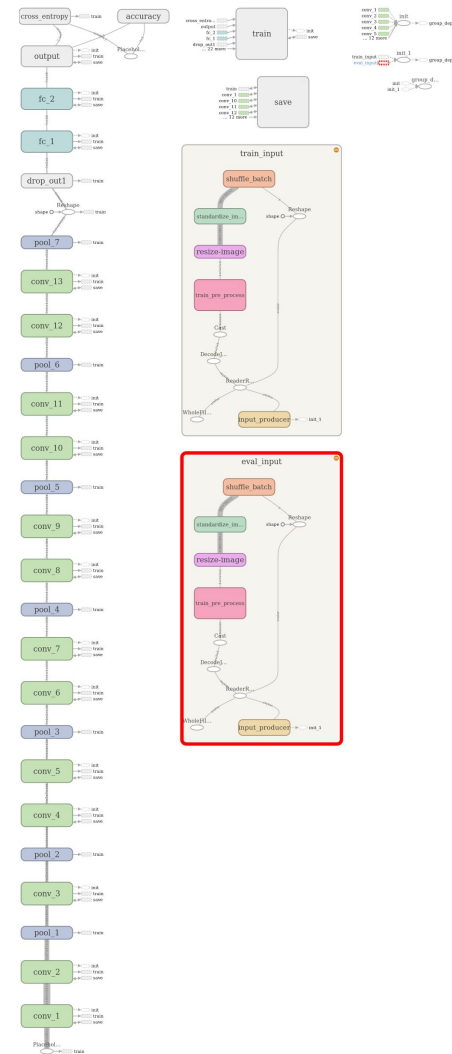
Tensorboard for
Visualization

Model Checkpointing

TF-slim, Keras,
TFlearn

Operations

- Design the Architecture
- Choose the cost function.
- Choose the Gradient to minimize the cost.
- Choose the metrics to calculate (optional).
- Define the placeholders
- Launch the graph and feed data to network



Design the Architecture using TensorFlow

```
#define a Architecture - 2 layer neural network
inputs = ...
with tf.name_scope("input_layer"):
    #define the weights
    w1= tf.Variable(tf.random_normal([784, 256]), name = "weights")
    #define the bias
    b1 = tf.Variable(tf.random_normal([256]), name = "bias")
    #multiply inputs with weights
    matmul = tf.matmul(inputs,w1)
    #Add bias to the matmul operation
    layer_1 = tf.add(matmul, b1)
    # Apply the activation function
    layer_1 = tf.nn.relu(layer_1)
```

Design the Architecture using TensorFlow

```
# 2nd layer
with tf.name_scope("hidden_layer"):
    w2 = tf.Variable(tf.random_normal([256, 128]), name = "weights")
    b2 = tf.Variable(tf.random_normal([128]), name = "bias")
    layer_2 = tf.add(tf.matmul(layer_1, w2), b2)
    layer_2 = tf.nn.relu(layer_2)

with tf.name_scope("output_layer"):
    w_o = tf.Variable(tf.random_normal([128, 10]), name = "weights")
    b_o = tf.Variable(tf.random_normal([10]), name = "bias")
    out_layer = tf.matmul(layer_2, w_o) + b_o
```

Calculate the loss

```
# Calculate the cost
with tf.name_scope("cost"):
    cost = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits=out_layer, labels=y_true))
    tf.summary.scalar("loss", cost)
```

Optimize the cost, Calculate the Accuracy

```
# calculate the gradient and back-prop the errors
with tf.name_scope("optimizer"):
    optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)

# Calculate the Accuracy:
with tf.name_scope("prediction"):
    prediction = tf.equal(tf.argmax(pred, 1), tf.argmax(y_true, 1))

with tf.name_scope("accuracy"):
    accuracy = tf.reduce_mean(tf.cast(prediction, tf.float32))
    tf.summary.scalar("accuracy", accuracy)
```


—

Initialize all the variables, Merge all the metrics calculated.

```
#Initialize the variables  
init = tf.global_variable_initializer()  
  
#merge all the summary operations to monitor metrics  
merged_summary_op = tf.summary.merge_all()
```

Launch the graph

```
# How initialize running ? Launch a session  
with tf.Session() as sess:  
    sess.run(init)
```

How to feed data to network ?

```
# How to feed data to the network ?
inputs = tf.placeholder(tf.float32, [None, 784]) # None so that we can decide on the batch size later
y_true = tf.placeholder(tf.float32, [None, 10])
```

Feed Data to Network

```
# How to feed input placeholders ?  
sess.run([optimizer], feed_dict = {inputs: train_features, y_true: train_labels})  
  
# How to feed batches of data ?  
for batch in range(n_batches):  
    sess.run([optimizer], feed_dict = {inputs: train_features[batch], y_true: train_labels[batch]})
```

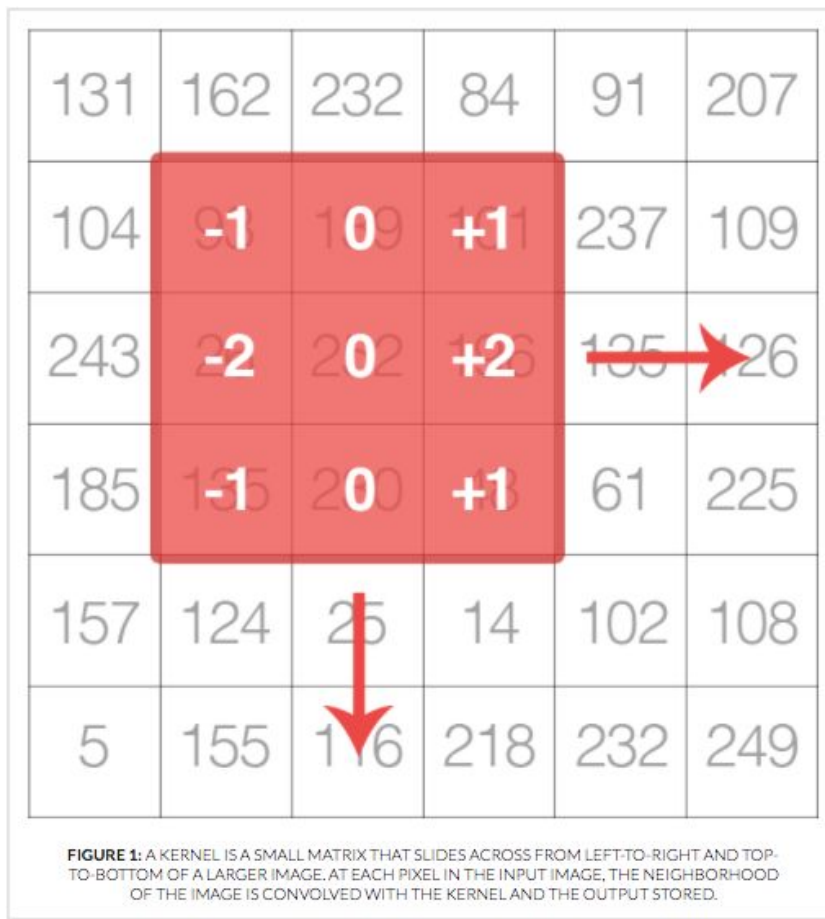
Feed Forward Neural Networks using Tensorflow

Logistic Regression using TensorFlow

A little deeper into Computer Vision

What are kernels?


- Kernels or convolutional matrix is a tiny matrix that is used for blurring, sharpening and edge detection.
- Convolution Process:
 - An input image
 - A kernel matrix that we are going to apply to the input image.
 - An output image to store the output of the input image convolved with the kernel.





Terminology in Convolution

- Strides
- Padding
- Filter maps

7 x 7 Input Volume

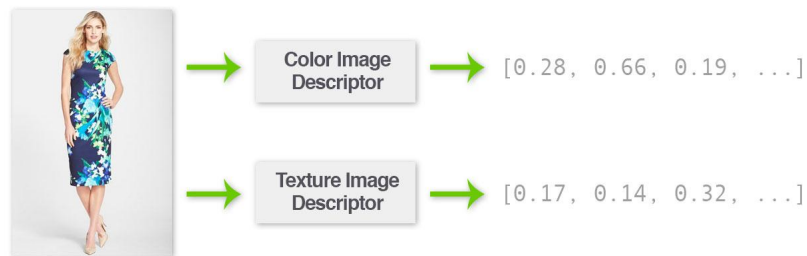


3 x 3 Output Volume





The input volume is 32 x 32 x 3. If we imagine two borders of zeros around the volume, this gives us a 36 x 36 x 3 volume. Then, when we apply our conv layer with our three 5 x 5 x 3 filters and a stride of 1, then we will also get a 32 x 32 x 3 output volume.

How to Vectorize an Image?



- An abstraction of an image used to characterize and numerically quantify the contents of an image. Normally real, integer, or binary valued. Simply put, a feature vector is a list of numbers used to represent an image.
- Color: Make a histogram with n-bins using a pixels. Normalize it and it becomes your color Image Descriptor
- Local Binary Patterns : Used to detect the patterns in Images locally.
- Most popularly used - Histogram of oriented Gradients.



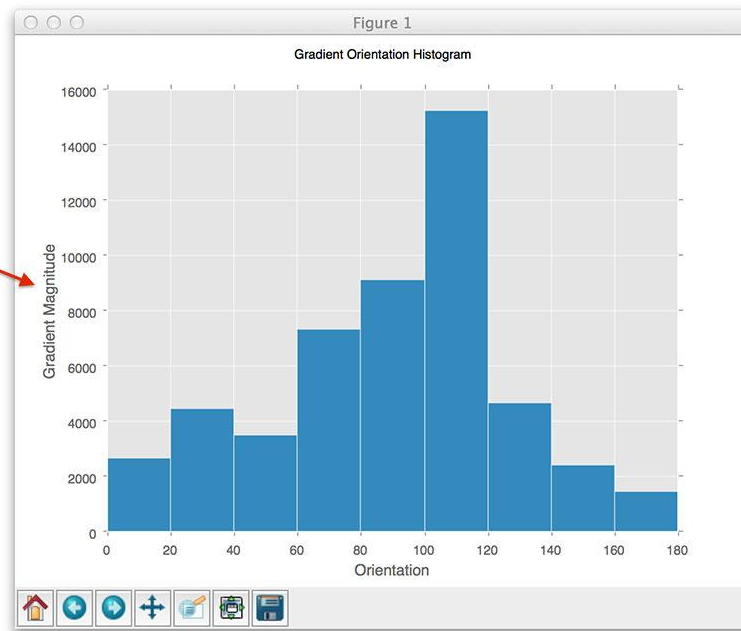

How to do Classification in Images using Image Descriptors?

- Histogram of Oriented Gradients.
 - Pixels per cell
 - Orientations
 - Cells_per_block

Example :




- 32 * 32 input Image
- Pixels per cell - $3 * 3 = 9$
- Cells per block = $2 * 2$
- Orientation = 9

Total Features = ?



How to do Classification in Images using Feature Descriptors?

- An Algorithm to detect Key points - example : FAST
- $R = 3$ Generally and we only consider points along the circle
- Now use SIFT Algorithm to detect features.
- FAST - detect edges, blobs etc
- SIFT extract features from this edge.
- We extract a patch for every key point and apply HOG to get a feature vector.
- RootSIFT: Take square root of each of your element.



16 x 16



Building an Image Classifier.

Extract Key Points and Feature Descriptors

N*M Dimensional vector for each Image

BAG of Visual Words,
Do clustering.

Define Number of clusters.
Domain specific

Multiple Features to
Single Histogram

Create a zero vector with
length as number of visual
words

For each Feature Descriptor.
Find nearest cluster and
update vector with +1

Resultant vector with count of
occurrence of each visual word
is our histogram.

CHOOSE YOUR MACHINE LEARNING
CLASSIFIER

Can We build a Image Search Engine In this Process?


Where does Neural Networks
comes into Picture?

Convolutional Neural Networks

Space Invariant Artificial Neural Network.

Inspired by How Human Perceive Object.

The First paper on CNNs is dating back to 1998. The Paper is called LeNet published by Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner.



IMAGENET

Early Problems

Not Enough
Computational
Power

Not Enough Data

Back Propagation
was not deep
enough

Important Terminology

Convolution Layer - Strides , Filters, Padding

Pooling Layers - Max , Avg, Fractional

Dropout , DropConnect

Activation Function

Loss Metrics





Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.



Activation Functions

Sigmoid

Relu

Leaky relu

Tanh

Maxout

Softplus

softmax

Data Pipelines in TensorFlow

Read the Image Locations

```
with tf.name_scope("train_input"):
    filenames = tf.train.match_filenames_once("/data/location/*.png", name = "filenamesholder")

    #create a queue of filenames
    filename_queue = tf.train.string_input_producer(filenames, num_epochs = num_epochs, shuffle = True)
```

Data Pipelines in TensorFlow

Decode the Image

```
with tf.name_scope("read_files"):
    # initiate a file reader
    reader = tf.WholeFileReader()

    # read a file
    key, record_string = reader.read(filename_queue)

    # decode the png and convert into float
    image = tf.image.decode_png(record_string, channels = 3)
    image = tf.cast(image, tf.float32)

with tf.name_scope("pre_process_image"):
    #Apply the pre_processing
    image = pre_process(image) # preprocess using tensorflow functions
```

Data Pipelines in TensorFlow

Enqueue the images and feed to the network batch_wise

```
with tf.name_scope("pre_process_image"):
    #Apply the pre_processing
    image = pre_process(image) # preprocess using tensorflow functions

# Queue the images and feed the network with a batch of images
with tf.name_scope("queue_images"):
    #Establish a queue (An Image storer, which stores images and feed to the network when required)
    min_after_dequeue = 100 # minimum number of images in the bin after dequeue
    capacity = min_after_dequeue + 3* batch_size # max number of images which can be stored in bin
    image.set_shape([36, 36, 3]) # set the shape of images in the queue
    key.set_shape([1]) # set the key(label) shape
    example_batch, label_batch = tf.train.shuffle_batch([example, label], batch_size = batch_size, num_thread =4,
                                                         capacity = capacity, min_after_dequeue = min_after_dequeue)
```

Data Pipelines in TensorFlow

How to Read Binary Files using tensorflow ?

```
with tf.name_scope("input_data"):
```

```
    label_bytes = 2 # CIFAR 100 has 0-99 labels #use 1 for CIFAR-10 (0-9)
```

```
    height, width, depth = 32, 32, 2
```

```
    image_bytes = 32* 32 * 3
```

```
record_bytes = label_bytes + image_bytes
```

```
reader = tf.FixedLengthRecordReader(record_bytes = record_bytes)
```

```
key, value = reader.read(filename_queue) #read filename queue.
```

```
record_bytes = tf.decode_raw(value, tf.uint8)
```

```
label = tf.cast(tf.strided_slice(record_bytes, [0], [label_bytes]), tf.int32)
```

```
depth_major = tf.reshape(tf.strided_slice(record_bytes, [label_bytes],  
                                          [label_bytes + image_bytes]), [depth, height, width])
```

```
uint8image = tf.transpose(depth_major, [1, 2, 0])
```

```
image = tf.cast(uint8image, tf.float32)
```

```
#Image and label to the file storage bin - tf.train.shuffle_batch
```

AlexNet

Image Input (224*224*3)

Kernel size = $11 * 11 * 3$, stride = 4 pixels

Output size? ($55 * 55$)

Conv_1 (96 filter Maps)

Local response normalization

Kernel size = $3 * 3$, stride = 2 pixels

Max Pooling

Kernel size = $5 * 5 * 96$, stride = 3 pixels

Output size? ($27 * 27$)

Conv_2 (256 filter Maps)

Kernel size = $3 * 3$, stride = 2 pixels

Output size? ($13 * 13$)

Max Pooling

Kernel size = $3 * 3 * 256$, stride = 1 pixels

Output size? ($13 * 13$)

Conv_3 (384 filter Maps)

Kernel size = $3 * 3 * 384$, stride = 1 pixels

Output size? ($13 * 13$)

Conv_4 (384 filter Maps)

Kernel size = $3 * 3 * 384$, stride = 1 pixels

Output size? ($13 * 13$)

Conv_5 (256 filter Maps)

Kernel size = $3 * 3$, stride = 2 pixels

Output size? ($6 * 6$)

Max Pooling

Total Features = $6 * 6 * 256 = 9216$

Fully Connected - 4096

Fully Connected - 4096

Output - 1000

**Neurons to train in Layer-1 =
 $((11 * 11 * 3) + 1) * 96 = 34,944$**

Fractional Max Pooling

General Pooling reduces the size of feature map by 1/4.

$$N_{in} = 2 * N_{out}$$

It Quickly reduces the size of feature map thereby not allowing to go deep and learn more features


$$N_{in} = 1.414 * N_{out} \text{ (Nearest integer)}$$

Random or Pseudo Random

Disjoint or Overlapping

What happens to the network when using FMP ?

11*11
C2
10*10
MP2
5*5
C2
4*4
MP2
2*2
C2
1*1



Design Architecture with less verbose code - code reusability

```
def conv2d(inputs, filters, kernel_size, strides, name, padding = "SAME"):
    shape = inputs.get_shape().as_list()
    with tf.name_scope(name):
        with tf.name_scope("weights"):
            weights = tf.Variable(tf.truncated_normal([kernel_size, kernel_size, shape[3], filters], stddev=0.1))
            variable_summaries(weights)
        with tf.name_scope("bias"):
            bias = tf.Variable(tf.zeros(shape = filters, dtype=tf.float32))
            variable_summaries(bias)
        nets = tf.nn.conv2d(inputs, weights, strides = [1, strides, strides, 1] , padding= padding)
        nets = tf.nn.bias_add(nets, bias)
        with tf.name_scope("leaky_relu"):
            nets = tf.maximum(nets, 0.1 * nets)
            variable_summaries(nets)
    return nets
```

Visualization using TensorBoard

```
def variable_summaries(var):  
    with tf.name_scope("summaries"):  
        mean = tf.reduce_mean(var)  
        tf.summary.scalar("mean", mean)  
        with tf.name_scope("stddev"):  
            stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))  
        tf.summary.scalar("stddev", stddev)  
        tf.summary.scalar("max", tf.reduce_max(var))  
        tf.summary.scalar("min", tf.reduce_min(var))  
        tf.summary.histogram("histogram", var)
```

Design Architecture - FMP

```
def fmp(inputs, pr = math.sqrt(2.), name = "pool", pseudo_random = False, overlapping = False):  
    with tf.name_scope(name):  
        nets = tf.nn.fractional_max_pool(inputs, pooling_ratio = [1, pr, pr, 1], pseudo_random = False, overlapping = False)  
    return nets[0]
```

Design Architecture - Fully connected

```
def fc(inputs, output_shape, name):
    shape = inputs.get_shape().as_list()
    with tf.name_scope(name):
        with tf.name_scope("weights"):
            weights= tf.Variable(tf.truncated_normal([shape[1], output_shape], stddev=0.1))
            variable_summaries(weights)
        with tf.name_scope("bias"):
            bias = tf.Variable(tf.zeros(shape = output_shape, dtype=tf.float32))
            variable_summaries(bias)
        nets = tf.add(tf.matmul(inputs, weights), bias)
    return nets
```

**Train the network with one
line of code.**

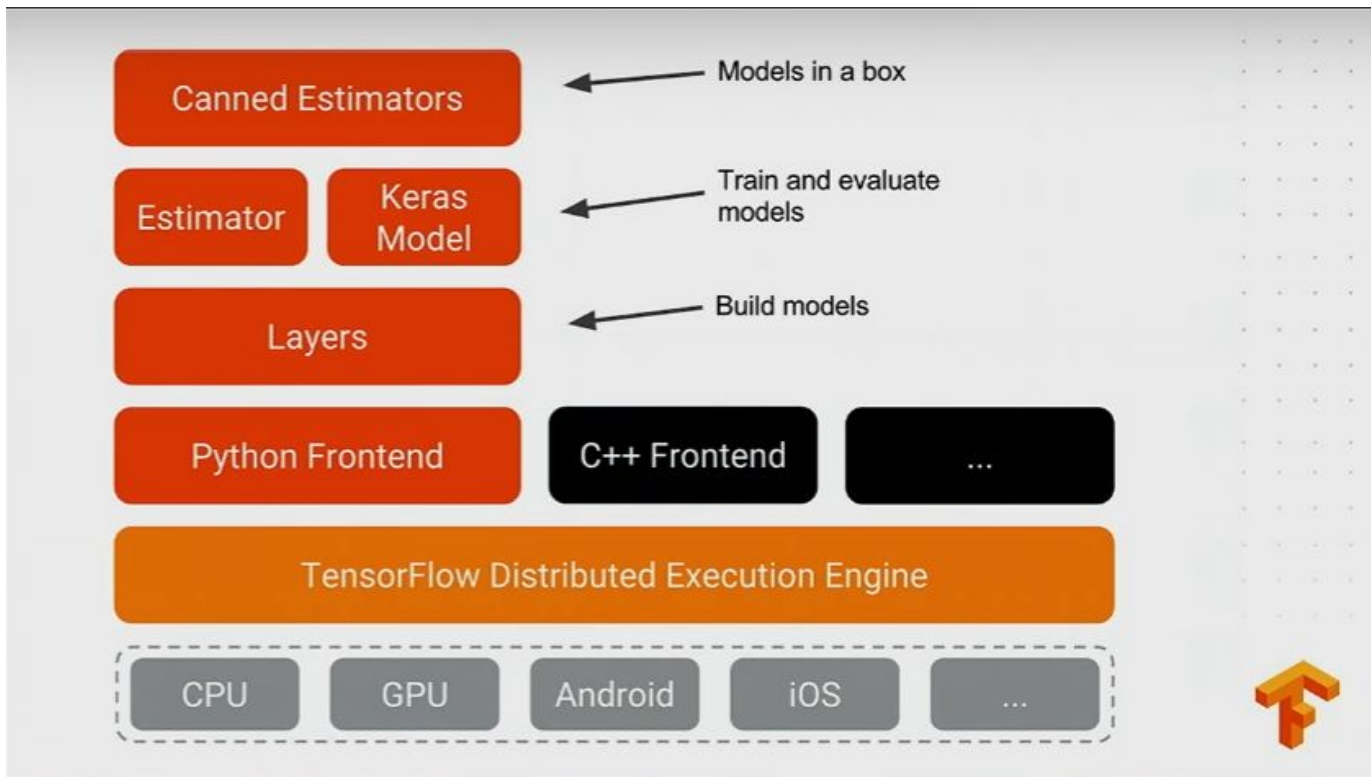
FMP on MNIST Dataset

FMP on CIFAR-100 Dataset

Dataset and the number of repeat tests	pseudorandom disjoint	random disjoint	pseudorandom overlapping	random overlapping
MNIST, 1 test	0.54	0.57	0.44	0.50
MNIST, 12 tests	0.38	0.37	0.34	0.32
CIFAR-100, 1 test	31.67	32.06	31.2	31.45
CIFAR-100, 12 tests	28.48	27.89	28.16	26.39

Table 1: MNIST and CIFAR-100 % test errors.

Everything under one hood



What we do at AthenaSight ...

Questions ?

Reach me at:

Email - vanapaliprakash@gmail.com

Twitter - @14prakash

—

Thank You